


JANASANKHYA · built

A synthetic population of India's informal workers, generated by a real CTGAN trained on IHDS-II microdata, enriched with cited structural priors and LLM-persona reasoning, and used to audit whether AI systems give worse welfare advice in low-resource languages.

At a glance




The architecture

every variable traces to real data or a citation

Four layers turn a real survey into an AI-fairness instrument.


JANASANKHYA — a synthetic population of India's informal workers, built to audit the AI systems that serve them



Layer 1 — the CTGAN is real, and it validates

Synthetic distributions are statistically close to the real IHDS-II data, and the model does not memorise individuals.


JANASANKHYA Layer 1 — CTGAN validation (real IHDS-II vs synthetic)



It preserves structural inequality, not just averages

The caste and gender wage gaps in the real data survive generation — this is what makes the population useful for fairness work.

CTGAN keeps the direction of structural inequality; it compresses the gender gap (real F:M=0.40 → synth 0.61) — which is what Layer 2 corrects



Eight synthetic workers


None is a real person; each is a statistically realistic informal worker with imputed AI-interaction variables.

state	sex	age	social_group	education	employment_type	annual_wage_inr	language	ai_usability_score	wage_below_minimum
Karnataka	Male	22	Dalit (SC)	Middle (6-9)	casual_daily	₹79,811	Kannada	2	True
Bihar	Male	37	OBC	Middle (6-9)	short_term_contract	₹115,079	Maithili	2	True
Himachal Pradesh	Male	44	Forward caste	Secondary (10-11)	casual_daily	₹92,966	Hindi	3	True
Uttar Pradesh	Female	20	Muslim	Middle (6-9)	casual_daily	₹35,813	Hindi	3	True
Kerala	Male	45	Christian/Sikh/Jain	Graduate (15)	regular_informal	₹149,675	Hindi	5	True

Maharashtra	Female	35	Forward caste	Middle (6-9)	regular_informal	₹33,368	Marathi	2	True
Uttar Pradesh	Male	21	Muslim	1-4 yrs	casual_daily	₹49,390	Hindi	1	True
Maharashtra	Female	48	Forward caste	1-4 yrs	casual_daily	₹175,365	Marathi	1	True

The finding — AI welfare advice degrades by language

Same synthetic worker, same welfare question, only the language changes. Recall = of the schemes the worker is actually entitled to, how many did the AI surface.



Query language	Mean welfare-scheme recall	Answer chars
English	0.75	1458
Hindi	0.56	1297
Bhojpuri	0.62	1343
Bengali	0.58	1348
Tamil	0.53	1281

Layers 2 + 3 — convergent validity

Two independent imputation methods agree: AI-usability scores from structural priors (Layer 2) and from LLM-persona reasoning (Layer 3) match within ±1 for **100%** of workers (correlation 0.062); rights-awareness mean abs error 0.086.

Data: IHDS-II individual file (ICPSR 36151, DS0001). Model: CTGAN (Xu et al., NeurIPS 2019), trained 300 epochs on 45,424 real informal workers. AI audit run against the local `claude` model. Structural priors from Jan Sahas (2020), ILO India Employment Report (2024), NITI Aayog (2022), GSMA (2024) — figures editable and to be pinned to primary sources during validation. Generated by the JANASANKHYA pipeline.